



```
In [1]: import os
import pandas as pd
import numpy as np
import cv2
import matplotlib.pyplot as plt
import plotly.express as px
import plotly.graph_objects as go
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix, accuracy_
import torch
import timm
import torchvision.transforms as transforms
from torch.utils.data import Dataset, DataLoader
import torch.nn as nn
import torch.optim as optim
from PIL import Image
```

```
In [2]: # Paths to datasets
train_dir = '/kaggle/input/retinal-disease-classification/Training_Set/Trainin
val_dir = '/kaggle/input/retinal-disease-classification/Evaluation_Set/Evaluat
test_dir = '/kaggle/input/retinal-disease-classification/Test_Set/Test_Set/Tes

train_labels_path = '/kaggle/input/retinal-disease-classification/Training_Set/
val_labels_path = '/kaggle/input/retinal-disease-classification/Evaluation_Set/
test_labels_path = '/kaggle/input/retinal-disease-classification/Test_Set/Test_
```

```
In [3]: # Load labels
train_labels = pd.read_csv(train_labels_path)
val_labels = pd.read_csv(val_labels_path)
test_labels = pd.read_csv(test_labels_path)
```

```
In [4]: # Print dataset shapes
print("Train labels shape:", train_labels.shape)
print("Validation labels shape:", val_labels.shape)
print("Test labels shape:", test_labels.shape)
```

```
Train labels shape: (1920, 47)
Validation labels shape: (640, 47)
Test labels shape: (640, 47)
```

```
In [5]: # Data Preprocessing
transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.RandomHorizontalFlip(p=0.5),
    transforms.RandomRotation(degrees=15),
    transforms.ColorJitter(brightness=0.2, contrast=0.2, saturation=0.2, hue=0),
    transforms.RandomAffine(degrees=0, translate=(0.1, 0.1)), # Small shifts
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
])
```

```
In [6]: class RetinalDataset(Dataset):
```

```

def __init__(self, image_dir, labels_df, transform=None):
    self.image_dir = image_dir
    self.labels_df = labels_df
    self.transform = transform

def __len__(self):
    return len(self.labels_df)

def __getitem__(self, idx):
    img_path = os.path.join(self.image_dir, f"{self.labels_df.iloc[idx]['Image']}")
    image = Image.open(img_path).convert('RGB')
    label = self.labels_df.iloc[idx]['Disease_Risk']
    if self.transform:
        image = self.transform(image)
    return image, torch.tensor(label, dtype=torch.float32)

```

In [7]: # Create datasets and loaders

```

train_dataset = RetinalDataset(train_dir, train_labels, transform=transform)
val_dataset = RetinalDataset(val_dir, val_labels, transform=transform)
test_dataset = RetinalDataset(test_dir, test_labels, transform=transform)

train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True)
val_loader = DataLoader(val_dataset, batch_size=32, shuffle=False)
test_loader = DataLoader(test_dataset, batch_size=32, shuffle=False)

```

In [8]:

```

import torch
print("CUDA Available:", torch.cuda.is_available())
print("Device:", torch.device("cuda" if torch.cuda.is_available() else "cpu"))

```

CUDA Available: True  
Device: cuda

In [9]: # Load pre-trained Vision Transformer (ViT) and Swin Transformer models

```

vit_model = timm.create_model('beit_base_patch16_224', pretrained=True, num_classes=2)
swin_model = timm.create_model('swin_base_patch4_window7_224', pretrained=True)

model.safetensors: 0% | 0.00/350M [00:00<?, ?B/s]
model.safetensors: 0% | 0.00/353M [00:00<?, ?B/s]

```

In [10]:

```

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
vit_model = vit_model.to(device)
swin_model = swin_model.to(device)

```

In [11]: # Loss function and optimizer

```

criterion = nn.BCEWithLogitsLoss()
vit_optimizer = optim.Adam(vit_model.parameters(), lr=1e-4, weight_decay=1e-5)
swin_optimizer = optim.Adam(swin_model.parameters(), lr=1e-4, weight_decay=1e-5)

```

In [12]:

```

def train_model(model, train_loader, val_loader, criterion, optimizer, scheduler):
    device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
    model.to(device)

    best_val_loss = float('inf')

```

```

patience_counter = 0

# Store losses and accuracies
train_losses, val_losses = [], []
train_accs, val_accs = [], []

for epoch in range(epochs):
    # Training Phase
    model.train()
    train_loss, correct, total = 0.0, 0, 0

    for images, labels in train_loader:
        images, labels = images.to(device), labels.to(device).unsqueeze(1)

        optimizer.zero_grad()
        outputs = model(images)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        train_loss += loss.item()
        preds = (torch.sigmoid(outputs) > 0.5).float()
        correct += (preds == labels).sum().item()
        total += labels.size(0)

    train_acc = correct / total
    train_loss /= len(train_loader)

    # Validation Phase
    model.eval()
    val_loss, val_correct, val_total = 0.0, 0, 0

    with torch.no_grad():
        for images, labels in val_loader:
            images, labels = images.to(device), labels.to(device).unsqueeze(1)
            outputs = model(images)
            loss = criterion(outputs, labels)
            val_loss += loss.item()

            preds = (torch.sigmoid(outputs) > 0.5).float()
            val_correct += (preds == labels).sum().item()
            val_total += labels.size(0)

    val_acc = val_correct / val_total
    val_loss /= len(val_loader)

# Store metrics
train_losses.append(train_loss)
val_losses.append(val_loss)
train_accs.append(train_acc)
val_accs.append(val_acc)

# Print Progress

```

```

print(f"Epoch {epoch+1}/{epochs} | Train Loss: {train_loss:.4f}, Train Acc: {train_acc:.4f}, Val Loss: {val_loss:.4f}, Val Acc: {val_acc:.4f}")

# Check for Best Model (Early Stopping)
if val_loss < best_val_loss:
    best_val_loss = val_loss
    torch.save(model.state_dict(), "best_model.pth")
    patience_counter = 0
else:
    patience_counter += 1

if patience_counter >= patience:
    print("Early stopping triggered! Stopping training.")
    break

# Adjust Learning Rate (if scheduler provided)
if scheduler:
    scheduler.step()

print("Training complete. Best Validation Loss:", best_val_loss)

return train_losses, val_losses, train_accs, val_accs

```

```

In [13]: # Loss function
criterion = nn.BCEWithLogitsLoss()

# Optimizers
vit_optimizer = optim.Adam(vit_model.parameters(), lr=1e-4, weight_decay=1e-5)
swin_optimizer = optim.Adam(swin_model.parameters(), lr=1e-4, weight_decay=1e-5)

# Learning Rate Scheduler (Reduce LR after 5 epochs)
vit_scheduler = torch.optim.lr_scheduler.StepLR(vit_optimizer, step_size=5, gamma=0.1)
swin_scheduler = torch.optim.lr_scheduler.StepLR(swin_optimizer, step_size=5, gamma=0.1)

# Train ViT
print("Training ViT Model:")
vit_train_losses, vit_val_losses, vit_train_accs, vit_val_accs = train_model(
    vit_model, train_loader, val_loader, criterion, vit_optimizer, vit_scheduler
)

# Train Swin Transformer
print("\nTraining Swin Transformer Model:")
swin_train_losses, swin_val_losses, swin_train_accs, swin_val_accs = train_model(
    swin_model, train_loader, val_loader, criterion, swin_optimizer, swin_scheduler
)

```

Training ViT Model:

Epoch 1/15 | Train Loss: 0.5243, Train Acc: 0.7906 | Val Loss: 0.5174, Val Acc: 0.7906  
Epoch 2/15 | Train Loss: 0.4999, Train Acc: 0.7911 | Val Loss: 0.5138, Val Acc: 0.7906  
Epoch 3/15 | Train Loss: 0.4936, Train Acc: 0.7911 | Val Loss: 0.4853, Val Acc: 0.7906  
Epoch 4/15 | Train Loss: 0.4835, Train Acc: 0.7911 | Val Loss: 0.4699, Val Acc: 0.7906  
Epoch 5/15 | Train Loss: 0.4774, Train Acc: 0.7958 | Val Loss: 0.4939, Val Acc: 0.7797  
Epoch 6/15 | Train Loss: 0.4609, Train Acc: 0.8000 | Val Loss: 0.4746, Val Acc: 0.7766  
Epoch 7/15 | Train Loss: 0.4490, Train Acc: 0.8026 | Val Loss: 0.4512, Val Acc: 0.8109  
Epoch 8/15 | Train Loss: 0.4315, Train Acc: 0.8177 | Val Loss: 0.4600, Val Acc: 0.8047  
Epoch 9/15 | Train Loss: 0.4254, Train Acc: 0.8193 | Val Loss: 0.4257, Val Acc: 0.8078  
Epoch 10/15 | Train Loss: 0.4124, Train Acc: 0.8203 | Val Loss: 0.4004, Val Acc: 0.8234  
Epoch 11/15 | Train Loss: 0.3945, Train Acc: 0.8344 | Val Loss: 0.3920, Val Acc: 0.8297  
Epoch 12/15 | Train Loss: 0.3811, Train Acc: 0.8427 | Val Loss: 0.3974, Val Acc: 0.8172  
Epoch 13/15 | Train Loss: 0.3831, Train Acc: 0.8333 | Val Loss: 0.3694, Val Acc: 0.8438  
Epoch 14/15 | Train Loss: 0.3765, Train Acc: 0.8417 | Val Loss: 0.4193, Val Acc: 0.8141  
Epoch 15/15 | Train Loss: 0.3834, Train Acc: 0.8380 | Val Loss: 0.3800, Val Acc: 0.8406  
Training complete. Best Validation Loss: 0.3694283176213503

Training Swin Transformer Model:

Epoch 1/15 | Train Loss: 0.4216, Train Acc: 0.8130 | Val Loss: 0.3429, Val Acc: 0.8187  
Epoch 2/15 | Train Loss: 0.3235, Train Acc: 0.8531 | Val Loss: 0.2737, Val Acc: 0.8703  
Epoch 3/15 | Train Loss: 0.2823, Train Acc: 0.8604 | Val Loss: 0.2564, Val Acc: 0.8703  
Epoch 4/15 | Train Loss: 0.2519, Train Acc: 0.8745 | Val Loss: 0.3421, Val Acc: 0.8250  
Epoch 5/15 | Train Loss: 0.2485, Train Acc: 0.8953 | Val Loss: 0.2488, Val Acc: 0.8875  
Epoch 6/15 | Train Loss: 0.2171, Train Acc: 0.9021 | Val Loss: 0.2785, Val Acc: 0.8859  
Epoch 7/15 | Train Loss: 0.1908, Train Acc: 0.9109 | Val Loss: 0.3362, Val Acc: 0.8688  
Epoch 9/15 | Train Loss: 0.1786, Train Acc: 0.9203 | Val Loss: 0.2269, Val Acc: 0.8875  
Epoch 10/15 | Train Loss: 0.1819, Train Acc: 0.9146 | Val Loss: 0.2550, Val Acc: 0.8906  
Epoch 11/15 | Train Loss: 0.1516, Train Acc: 0.9318 | Val Loss: 0.2597, Val Acc: 0.8859

```
Epoch 12/15 | Train Loss: 0.1283, Train Acc: 0.9443 | Val Loss: 0.3079, Val Ac  
c: 0.8719  
Early stopping triggered! Stopping training.  
Training complete. Best Validation Loss: 0.22686027251183988
```

```
In [19]: # Evaluation function  
def evaluate_model(model, test_loader, model_name):  
    device = torch.device("cuda" if torch.cuda.is_available() else "cpu")  
    model.eval()  
    all_preds, all_labels = [], []  
  
    with torch.no_grad():  
        for images, labels in test_loader:  
            images, labels = images.to(device), labels.to(device).unsqueeze(1)  
            outputs = model(images)  
            preds = (torch.sigmoid(outputs) > 0.5).float()  
            all_preds.extend(preds.cpu().numpy())  
            all_labels.extend(labels.cpu().numpy())  
  
    acc = accuracy_score(all_labels, all_preds)  
    print(f"{model_name} Test Accuracy: {acc:.4f}")  
    print(classification_report(all_labels, all_preds))  
    return acc
```

```
In [15]: # Evaluate both models  
vit_test_acc = evaluate_model(vit_model, test_loader, "ViT")  
swin_test_acc = evaluate_model(swin_model, test_loader, "Swin Transformer")
```

```
ViT Test Accuracy: 0.8422  
precision recall f1-score support  
0.0 0.60 0.75 0.67 134  
1.0 0.93 0.87 0.90 506  
  
accuracy 0.84 640  
macro avg 0.76 0.81 0.78 640  
weighted avg 0.86 0.84 0.85 640
```

```
Swin Transformer Test Accuracy: 0.9328  
precision recall f1-score support  
0.0 0.79 0.93 0.85 134  
1.0 0.98 0.93 0.96 506  
  
accuracy 0.93 640  
macro avg 0.88 0.93 0.90 640  
weighted avg 0.94 0.93 0.93 640
```

```
In [16]: import matplotlib.pyplot as plt  
  
# Debugging: Check the lengths of lists  
print("Length of training/validation loss & accuracy lists:")  
print("ViT Train Losses:", len(vit_train_losses))
```

```

print("ViT Val Losses:", len(vit_val_losses))
print("Swin Train Losses:", len(swin_train_losses))
print("Swin Val Losses:", len(swin_val_losses))

print("ViT Train Accs:", len(vit_train_accs))
print("ViT Val Accs:", len(vit_val_accs))
print("Swin Train Accs:", len(swin_train_accs))
print("Swin Val Accs:", len(swin_val_accs))

# Find the shortest epoch count (to avoid shape mismatch errors)
min_epochs = min(
    len(vit_train_losses), len(vit_val_losses),
    len(swin_train_losses), len(swin_val_losses),
    len(vit_train_accs), len(vit_val_accs),
    len(swin_train_accs), len(swin_val_accs)
)

# Trim all lists to match the shortest length
vit_train_losses = vit_train_losses[:min_epochs]
vit_val_losses = vit_val_losses[:min_epochs]
swin_train_losses = swin_train_losses[:min_epochs]
swin_val_losses = swin_val_losses[:min_epochs]

vit_train_accs = vit_train_accs[:min_epochs]
vit_val_accs = vit_val_accs[:min_epochs]
swin_train_accs = swin_train_accs[:min_epochs]
swin_val_accs = swin_val_accs[:min_epochs]

# Define the epoch range
epochs_range = range(min_epochs)

# Plot loss and accuracy comparison
plt.figure(figsize=(12, 5))

# Loss Plot
plt.subplot(1, 2, 1)
plt.plot(epochs_range, vit_train_losses, label="ViT Train Loss", marker="o")
plt.plot(epochs_range, vit_val_losses, label="ViT Val Loss", marker="s")
plt.plot(epochs_range, swin_train_losses, label="Swin Train Loss", marker="^")
plt.plot(epochs_range, swin_val_losses, label="Swin Val Loss", marker="d")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.title("Loss Comparison")

# Accuracy Plot
plt.subplot(1, 2, 2)
plt.plot(epochs_range, vit_train_accs, label="ViT Train Accuracy", marker="o")
plt.plot(epochs_range, vit_val_accs, label="ViT Val Accuracy", marker="s")
plt.plot(epochs_range, swin_train_accs, label="Swin Train Accuracy", marker="^")
plt.plot(epochs_range, swin_val_accs, label="Swin Val Accuracy", marker="d")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")

```

```

plt.legend()
plt.title("Accuracy Comparison")


plt.show()

# Print final test accuracy for both models
print(f"\nFinal Test Accuracy - ViT: {vit_test_acc:.4f}, Swin Transformer: {swin_test_acc:.4f}")

```

Length of training/validation loss & accuracy lists:

ViT Train Losses: 15  
 ViT Val Losses: 15  
 Swin Train Losses: 12  
 Swin Val Losses: 12  
 ViT Train Accs: 15  
 ViT Val Accs: 15  
 Swin Train Accs: 12  
 Swin Val Accs: 12



Final Test Accuracy - ViT: 0.8422, Swin Transformer: 0.9328

```

In [17]: # Function to test the models on 10 images
def test_on_sample_images(models, test_dataset, model_names, num_samples=10):
    device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

    fig, axes = plt.subplots(2, num_samples, figsize=(20, 6))
    sample_indices = np.random.choice(len(test_dataset), num_samples, replace=False)

    for model, model_name, row in zip(models, model_names, range(2)):
        model.eval()
        with torch.no_grad():
            for i, idx in enumerate(sample_indices):
                image, label = test_dataset[idx]
                image = image.unsqueeze(0).to(device)
                output = model(image)
                pred = torch.sigmoid(output).item()
                predicted_label = "Diseased" if pred > 0.5 else "Healthy"

                image = image.squeeze(0).cpu().permute(1, 2, 0).numpy()

```

```


image = (image * [0.229, 0.224, 0.225]) + [0.485, 0.456, 0.406
image = np.clip(image, 0, 1)

ax = axes[row, i]
ax.imshow(image)
ax.set_title(f"{model_name}\nPred: {predicted_label}\nTrue: {true_label}")
ax.axis("off")

plt.tight_layout()
plt.show()

# Run the test function for both models
test_on_sample_images([vit_model, swin_model], test_dataset, ["ViT", "Swin Transformer"])

```



In [26]: `torch.save(vit_model.state_dict(), "beit1.pth")  
torch.save(swin_model.state_dict(), "swin1.pth")`